

Reinforcement Learning

Lecture 1

Lawrence Carin
Duke University



- ❑ Consider an MD, seeking to treat patients in an effective manner, while minimizing costs
- ❑ The state of health for each patient is observed in terms of a set of clinical variables
- ❑ Assume there are a set of actions the MD can perform
- ❑ Each action will impact the health state of the patient (possibly no change in the state)
- ❑ Seek a policy that achieves the best outcomes at the lowest average price
- ❑ Could apply to particular types of patients (e.g., diabetics) or in particular settings within a health system (e.g., operating room)



- ❑ Assume the set of observations defining the patient **state** represented by vector s
- ❑ Assume the set of **actions** that may be taken is $A = \{a_1, \dots, a_m\}$
- ❑ There is **randomness** in how a patient in state s will respond to a given action
- ❑ Let $P(s, a, s')$ represent the *probability* that when a patient is in state s and the MD takes action a , the patients state will change to s'
- ❑ Let $r(s, a, s')$ represent the “**reward**” associated with the MD taking action a for a patient in state s , and then the patient transiting to new state of health s'
- ❑ The reward $r(s, a, s')$ may reflect the *immediate* impact for the patient of the change $s \rightarrow s'$ and also the cost of action a

- ❑ The MD interacts with the patient through a series of actions, patient state changes, and rewards/costs

$$\dots s_{t-1} \ a_{t-1} \ r_{t-1} \ s_t \ a_t \ r_t \ s_{t+1} \ \dots$$



- ❑ Goal: Develop a **policy** that defines for the MD the optimal action a to take when presented by a patient in state s ; the policy may define *standard of care*
- ❑ The optimal policy will maximize the average reward over time, with reward accounting for patient outcome and costs
- ❑ The policy should be non-myopic, in that it thinks ahead, to the long-run impact of actions
- ❑ The policy will typically weight impacts in the near-term more highly than what happens in the long run

Big Challenge



- We typically do not know $P(s, a, s')$
- How can we learn a policy without this?
- We can just experience/try things, keep a record of outcomes, and adapt and adjust
- Reinforce actions for particular patient states that are rewarding
- Discourage actions that are expensive and yield poor outcomes
- In many ways this is how medicine works, over time

- Reinforcement learning is the formalization of this challenge
- Addresses sequential decision making in an uncertain (stochastic) world



Medicine/Health

- Reinforcement learning is the formalization of this challenge
- Addresses sequential decision making in an uncertain (stochastic) world



Medicine/Health



Monitoring/Maintenance of Factory

- Reinforcement learning is the formalization of this challenge
- Addresses sequential decision making in an uncertain (stochastic) world



Medicine/Health



Monitoring/Maintenance of Factory



Investing

- Reinforcement learning is the formalization of this challenge
- Addresses sequential decision making in an uncertain (stochastic) world



Medicine/Health



Monitoring/Maintenance of Factory



Investing



Illustrative Example



☐ Consider diabetes control

☐ Assume that the patient **state** s is defined by the **minimum and maximum glucose concentration from previous day**

☐ The **action** a may be the **rate of continuous insulin supply** and the **bolus dose**

Illustrative Example



- ❑ Consider diabetes control
- ❑ Assume that the patient **state** s is defined by the **minimum and maximum glucose concentration from previous day**
- ❑ The **action** a may be the **rate of continuous insulin supply** and the **bolus dose**
- ❑ $r(s, a, s')$ may be defined to reward (punish) glucose levels that are desirable (undesirable)
- ❑ Assume that we may specify $r(s, a, s')$

Solution Setup



❑ Assume that the patient **state** s is defined by the **minimum and maximum glucose concentration from previous day**

- Discretize the continuous range of the state values into n bins

❑ The **action** a may be the **rate of continuous insulin supply** and the **bolus dose**

- Discretize the continuous range of the action values into m bins

Solution Setup



❑ Assume that the patient **state** s is defined by the **minimum and maximum glucose concentration from previous day**

- Discretize the continuous range of the state values into n bins

❑ The **action** a may be the **rate of continuous insulin supply** and the **bolus dose**

- Discretize the continuous range of the action values into m bins

❑ The Q function $Q(s, a)$ is an $n \times m$ matrix, denoting the value of taking action a when in state s

Simple, Intuitive Solution



❑ Set initial values of $n \times m$ matrix $Q(s, a)$ based on prior available medical knowledge, or set at random

❑ After initializing $Q(s, a)$, take an action a when patient is in particular state s , and then observe the new state s'

$$(s, a) \rightarrow s', \quad r(s, a, s')$$

Simple, Intuitive Solution



❑ Set initial values of $n \times m$ matrix $Q(s, a)$ based on prior available medical knowledge, or set at random

❑ After initializing $Q(s, a)$, take an action a when patient is in particular state s , and then observe the new state s'

$$(s, a) \rightarrow s', \quad r(s, a, s')$$

❑ Assume that the prior/old value for the Q function when taking action a in state s is $Q^{old}(s, a)$

Simple, Intuitive Solution



- ❑ Set initial values of $n \times m$ matrix $Q(s, a)$ based on prior available medical knowledge, or set at random

- ❑ After initializing $Q(s, a)$, take an action a when patient is in particular state s , and then observe the new state s'

$$(s, a) \rightarrow s', \quad r(s, a, s')$$

- ❑ Assume that the prior/old value for the Q function when taking action a in state s is $Q^{old}(s, a)$

- ❑ We may consider the update rule:

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0,1)$$

Simple, Intuitive Solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0,1)$$



Temporal Difference (TD)

- ❑ α is called the “learning rate,” and controls the relative balance between our old estimate $Q^{old}(s, a)$ and new information provided by $r(s, a, s')$
- ❑ If the TD is positive, the inferred value of taking action a in state s is increased; if TD negative, the value of taking action a in state s is diminished

Simple, Intuitive Solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0,1)$$



- ❑ If the reward $r(s, a, s')$ is large (small) then $Q^{new}(s, a)$ is typically increased (diminished)

Simple, Intuitive Solution



$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0,1)$$

- ❑ If the reward $r(s, a, s')$ is large (small) then $Q^{new}(s, a)$ is typically increased (diminished)
- ❑ Problem: This only accounts for the immediate reward $r(s, a, s')$
- ❑ Doesn't account for what may happen subsequently once the patient state changes to s'

Example: Problem With Myopic Policy



$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0,1)$$

- ❑ Assume state of patient s corresponds to severe poor health
- ❑ A particular action a may have probable positive immediate reward $r(s, a, s')$
- ❑ However, there may be serious long-term complications (e.g., loss of opportunity to have children)
- ❑ A policy that only accounts for the immediate reward would not account for long-term, later consequences

Extension

- Recall our simple solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0, 1)$$



Extension



- ❑ Recall our simple solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0, 1)$$

- ❑ Consider the extension

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') + \gamma \cdot \max_{a'} Q^{old}(s', a') - Q^{old}(s, a)]$$

Extension



- ❑ Recall our simple solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0, 1)$$

- ❑ Consider the extension

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') + \gamma \cdot \max_{a'} Q^{old}(s', a') - Q^{old}(s, a)]$$



Extension



- ❑ Recall our simple solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0, 1)$$

- ❑ Consider the extension

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [\underbrace{r(s, a, s')}_{\text{Immediate Reward}} + \underbrace{\gamma \cdot \max_{a'} Q^{old}(s', a')}_{\substack{\text{Discount Factor Between 0 and 1} \\ \text{Expected Future Rewards} \\ \text{Based on Optimal} \\ \text{Policy}}} - Q^{old}(s, a)]$$

Extension



- ❑ Recall our simple solution

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') - Q^{old}(s, a)], \text{ with } \alpha \in (0, 1)$$

- ❑ How about extending it as

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + \alpha \cdot [r(s, a, s') + \underbrace{\gamma \cdot \max_{a'} Q^{old}(s', a') - Q^{old}(s, a)}_{\text{Non-Myopic Temporal Difference (TD)}}]$$

Non-Myopic Temporal Difference (TD)

Q Learning



$$Q^{new}(s, a) \leftarrow (1 - \alpha) \cdot Q^{old}(s, a) + \alpha \cdot [r(s, a, s') + \gamma \cdot \max_{a'} Q^{old}(s', a')]$$

- ❑ Based on simple logic and intuition, we have derived an algorithm for a system/MD to learn based on experience
- ❑ This is actually a widely employed method for reinforcement learning, called Q Learning

Q Learning



$$Q^{new}(s, a) \leftarrow (1 - \alpha) \cdot Q^{old}(s, a) + \alpha \cdot [r(s, a, s') + \gamma \cdot \max_{a'} Q^{old}(s', a')]$$

- ❑ Based on simple logic and intuition, we have derived an algorithm for a system/MD to learn based on experience
- ❑ This is actually a widely employed method for reinforcement learning, called Q Learning
- ❑ Based on the learned matrix $Q(s, a)$ the policy that is typically employed is


$$\pi(a; s) = \operatorname{argmax}_a Q(s, a)$$

Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

Q Learning

Discount factor, $\gamma \in [0,1]$


$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$



Learning rate, $\alpha \in [0,1]$

Q Learning


$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$



Previous Estimate of Value of
Action a_t when in state s_t

Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$



Updated Estimate of Value of
Action a_t when in state s_t
After Observing r_t

Q Learning

The diagram illustrates the Q-learning update equation with several annotations:

- Learning rate, $\alpha \in [0,1]$** : A red label with a blue bracket pointing to the $(1 - \alpha)$ term and the α multiplier.
- Discount factor, $\gamma \in [0,1]$** : A red label with a blue arrow pointing to the γ term.
- Previous Estimate of Value of Action a_t when in state s_t** : A red label with a blue bracket pointing to the $Q^{old}(s_t, a_t)$ term.
- Updated Estimate of Value of Action a_t when in state s_t After Observing r_t** : A red label with a blue bracket pointing to the entire right-hand side of the equation, $(1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$.

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

❑ Don't need to actually perform the action defined by $\operatorname{argmax}_a Q^{old}(s_{t+1}, a)$

Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

❑ Don't need to actually perform the action defined by $\operatorname{argmax}_a Q^{old}(s_{t+1}, a)$

❑ Next action a_{t+1} may be based on $\operatorname{argmax}_{a_{t+1}} Q^{new}(s_{t+1}, a_{t+1})$, we then observe immediate reward r_{t+1} and new state s_{t+2}

... s_t a_t r_t s_{t+1} a_{t+1} r_{t+1} s_{t+2} ...

Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

❑ Don't need to actually perform the action defined by $\arg\max_a Q^{old}(s_{t+1}, a)$

❑ Next action a_{t+1} may be based on $\arg\max_{a_{t+1}} Q^{new}(s_{t+1}, a_{t+1})$, we then observe immediate reward r_{t+1} and new state s_{t+2}

... s_t a_t r_t s_{t+1} a_{t+1} r_{t+1} s_{t+2} ...

❑ Sequentially and continually update the Q function $Q^{new}(s_t, a_t)$, which is an $n \times m$ matrix

Learning Rate

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$



Learning rate, $\alpha \in [0,1]$

- ❑ Convergence guaranteed after sufficient experience, if learning rate diminishes with time
- ❑ In practice one may set $\alpha = 0.1$
- ❑ This online Q-learning setup allows one to learn an optimal policy based directly on experience

Q Learning vs. SARSA

□ **Q-learning**, from previous slides:

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

Assumes access to s_t a_t r_t s_{t+1} (SARS); we don't implement a_{t+1} to update the Q functions

Q Learning vs. SARSA

- ❑ **Q-learning**, from previous slides:

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

Assumes access to s_t a_t r_t s_{t+1} (SARS); we don't implement a_{t+1} to update the Q functions

- ❑ **SARSA** based learning considers the slightly modified

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot Q^{old}(s_{t+1}, a_{t+1})]$$

Assumes access to s_t a_t r_t s_{t+1} a_{t+1} (SARSA); a_{t+1} may be specified by some chosen mechanism

ϵ -Greedy: Exploration and Exploitation

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot Q^{old}(s_{t+1}, a_{t+1})]$$

□ Assume we observe s_t a_t r_t s_{t+1}

ϵ -Greedy: Exploration and Exploitation

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot Q^{old}(s_{t+1}, a_{t+1})]$$

☐ Assume we observe s_t a_t r_t s_{t+1}

☐ Which next action a_{t+1} should we take within Q learning or SARSA?

ϵ -Greedy: Exploration and Exploitation

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot Q^{old}(s_{t+1}, a_{t+1})]$$

□ Assume we observe $s_t \ a_t \ r_t \ s_{t+1}$

□ Which next action a_{t+1} should we take within Q learning or SARSA?

□ ϵ -Greedy (small ϵ):

- With probability ϵ , choose a_{t+1} at random
- With probability $1 - \epsilon$, next action $a_{t+1} = \underset{a}{\operatorname{argmax}} Q^{old}(s_{t+1}, a)$

ϵ -Greedy: Exploration and Exploitation

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot Q^{old}(s_{t+1}, a_{t+1})]$$

❑ Assume we observe s_t a_t r_t s_{t+1}

❑ Which next action a_{t+1} should we take within Q learning or SARSA?

❑ ϵ -Greedy (small ϵ):

- With probability ϵ , choose a_{t+1} at random
- With probability $1 - \epsilon$, next action $a_{t+1} = \underset{a}{\operatorname{argmax}} Q^{old}(s_{t+1}, a)$



Allows exploration with probability ϵ

Reinforcement Learning

Lecture 2

Lawrence Carin
Duke University


Limitations of Tabular Q Learning

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot \max_a Q^{old}(s_{t+1}, a)]$$

❑ Simple update rule, but

- The Q function is stored as a table/matrix, which is impractical when considering a large number of states and actions
- No real capacity to generalize across sequence types, because the tabular Q function does not have a functional form

Represent Q Function as a Neural Network



- ❑ The Q-function modeled via a neural network
- ❑ The state is input (e.g., an image with N pixels), and at the output we model the Q-function $Q(s, a)$ for each possible action
- ❑ Represent the neural network model as $Q(s, a; \theta)$ where θ represent the neural network parameters we wish to learn
- ❑ After we learn θ the policy is


$$\pi(s) = \operatorname{argmax}_a Q(s, a; \theta)$$

Deep Q Learning (DQN)

- If we update θ like in Q learning, with every new observed s_t, a_t, r_t, s'_{t+1} , at each step we seek to minimize

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$

- To move θ such that $U(\theta; s_t, a_t)$ is made smaller, we update model parameters as




Gradient Descent

- If we update θ like in Q learning, with every new observed s_t, a_t, r_t, s'_{t+1} , at each step we seek to minimize

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$

- To move θ such that $U(\theta; s_t, a_t)$ is made smaller, we update model parameters as

$$\theta^{new} \leftarrow \theta^{old} - \alpha \cdot \nabla_{\theta} U(\theta; s_t, a_t)|_{\theta=\theta^{old}}$$



Gradient Descent


- If we update θ like in Q learning, with every new observed s_t, a_t, r_t, s'_{t+1} , at each step we seek to minimize

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$

- To move θ such that $U(\theta; s_t, a_t)$ is made smaller, we update model parameters as

$$\theta^{new} \leftarrow \theta^{old} - \alpha \cdot \nabla_{\theta} U(\theta; s_t, a_t)|_{\theta=\theta^{old}}$$

Multi-dimensional “slope”



Gradient Descent

□ If we update θ like in Q learning, with every new observed s_t, a_t, r_t, s'_{t+1} , at each step we seek to minimize

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$

□ To move θ such that $U(\theta; s_t, a_t)$ is made smaller, we update model parameters as

$$\theta^{new} \leftarrow \theta^{old} - \alpha \cdot \nabla_{\theta} U(\theta; s_t, a_t)|_{\theta=\theta^{old}}$$

$$\theta^{new} \leftarrow \theta^{old} - \alpha \cdot \left[Q(s_t, a_t; \theta^{old}) - r_t - \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}}$$

Gradient Descent

$$\theta^{new} \leftarrow \theta^{old} + \alpha \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta) |_{\theta=\theta^{old}}$$



Looks a lot like the prior Q learning

Gradient Descent

$$\theta^{new} \leftarrow \theta^{old} + \alpha \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta) \big|_{\theta=\theta^{old}}$$



Modification to account for
neural network parameters θ

Gradient Descent

$$\theta^{new} \leftarrow \theta^{old} + \alpha \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta) |_{\theta=\theta^{old}}$$

Looks a lot like the prior Q learning

Modification to account for
neural network parameters θ

- ❑ One may show rigorously that this is directly related to the prior sequential Q learning
- ❑ Details not very difficult (first-order Taylor series expansion), but beyond the scope of our discussion (see **Appendix**)

Revisit What Deep Q Network Seeks to Learn

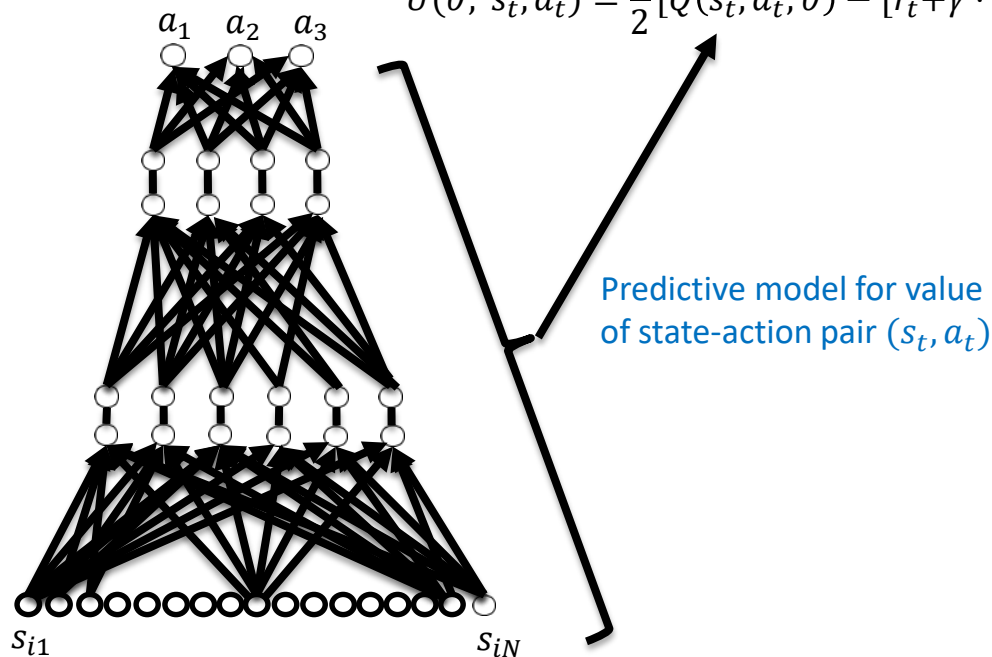
- Seek parameters θ that minimize prediction error

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$

Revisit What Deep Q Network Seeks to Learn


- Seek parameters θ that minimize prediction error

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$




Revisit What Deep Q Network Seeks to Learn

- Seek parameters θ that minimize prediction error



Revisit What Deep Q Network Seeks to Learn


- Seek parameters θ that minimize prediction error



Power of Deep Q Networks for RL With Images

- Seek parameters θ that minimize prediction error


$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$



Power of Deep Q Networks for RL With Images

- Seek parameters θ that minimize prediction error


$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$



Power of Deep Q Networks for RL With Images

- Seek parameters θ that minimize prediction error


$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$




Power of Deep Q Networks for RL With Images


- Seek parameters θ that minimize prediction error

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$





Atari Games



Final Thoughts

- ❑ Tabular Q Learning employs the update rule

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot Q^{old}(s_t, a_t) + \alpha \cdot [r_t + \gamma \cdot Q^{old}(s_{t+1}, a_{t+1})]$$

- ❑ Deep Q Learning uses a deep neural network with parameters θ to minimize the functional

$$U(\theta; s_t, a_t) = \frac{1}{2} [Q(s_t, a_t; \theta) - [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old})]]^2$$

- ❑ Reinforcement learning has been studied for decades, and there are many other methods that we have not considered here, and are worth learning about for those interested
- ❑ There is also a rich theoretical foundation to RL, and therefore much is known about the fundamentals of these methods

Appendix 1

Introduction to Markov Decision Process (MDP)

Markov Decision Process (MDP)

- ❑ A set of **states** $S = (s_1, \dots, s_n)$

We observe the state of the system/environment/subject

Markov Decision Process (MDP)

- ❑ A set of **states** $S = (s_1, \dots, s_n)$

We observe the state of the system/environment/subject

- ❑ A set of **actions** $A = (a_1, \dots, a_m)$

We may select an action to take, given that the system is in an observed state s

Markov Decision Process (MDP)

- ❑ A set of **states** $S = (s_1, \dots, s_n)$

We observe the state of the system/environment/subject

- ❑ A set of **actions** $A = (a_1, \dots, a_m)$

We may select an action to take, given that the system is in an observed state s

- ❑ **State-transition probabilities** $P(s, a, s')$

Defines the probability of transitioning $s \rightarrow s'$ upon taking action a when in state s ; $\sum_{s' \in S} P(s, a, s') = 1$

Markov Decision Process (MDP)

- ❑ A set of **states** $S = (s_1, \dots, s_n)$

We observe the state of the system/environment/subject

- ❑ A set of **actions** $A = (a_1, \dots, a_m)$

We may select an action to take, given that the system is in an observed state s

- ❑ **State-transition probabilities** $P(s, a, s')$

Defines the probability of transitioning $s \rightarrow s'$ upon taking action a when in state s ; $\sum_{s' \in S} P(s, a, s') = 1$

- ❑ Set of **rewards** $r(s, a, s')$ reflecting the *immediate* reward of taking action a in state s and transiting to state s'

Markov Decision Process (MDP)

- ❑ A set of **states** $S = (s_1, \dots, s_n)$

We observe the state of the system/environment/subject

- ❑ A set of **actions** $A = (a_1, \dots, a_m)$

We may select an action to take, given that the system is in an observed state s

- ❑ **State-transition probabilities** $P(s, a, s')$


Defines the probability of transitioning $s \rightarrow s'$ upon taking action a when in state s ; $\sum_{s' \in S} P(s, a, s') = 1$

- ❑ Set of **rewards** $r(s, a, s')$ reflecting the *immediate* reward of taking action a in state s and transiting to state s'

- ❑ **Discount factor** $\gamma \in [0, 1)$

Reflects the degree to which future rewards are discounted relative to immediate rewards

Markov Decision Process (MDP)



Policy

- ❑ A policy π specifies which action $a \in A$ is taken when in state $s \in S$
- ❑ The policy may be stochastic, with $\pi(a; s)$ reflecting the *probability* of action a when in state s ; $\sum_{i=1}^m \pi(a_i; s) = 1$

Policy

- ❑ A policy π specifies which action $a \in A$ is taken when in state $s \in S$
- ❑ The policy may be stochastic, with $\pi(a; s)$ reflecting the *probability* of action a when in state s ; $\sum_{i=1}^m \pi(a_i; s) = 1$
- ❑ Consider the sequence $(s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{N-1}, a_{N-1}, r_{N-1}, s_N, a_N, r_N)$
 - The total reward for this sequence is $r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^N r_N$
 - This total reward is a *random variable*, because the sequence of states visited is random
- ❑ We wish to learn the policy π to maximize the expected reward of the “agent” over a sequence of states, actions and rewards

Expected (Average) Reward of a Policy

- ❑ The state-action value function $Q^\pi(s, a)$ of any policy π indicates the expected (average) discounted total reward when taking action a in state s and following the optimal policy thereafter

$$Q^\pi(s, a) = E_{a_t \sim \pi; s_t \sim P}(\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a)$$

- ❑ We may re-express $Q^\pi(s, a)$ as

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a'} Q^\pi(s', a')$$

- ❑ Immediate *expected* reward $R(s, a) = \sum_{s'} P(s, a, s') r(s, a, s')$

Deterministic Policy

- ❑ Assumed the policy selects the single most rewarding action when in state s

$$\pi(s) = \arg \max_{a \in A} Q^\pi(s, a)$$

- ❑ Bellman's equation

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a' \in A} Q^\pi(s', a')$$

- ❑ Suggests an iterative solution to learn the optimal deterministic policy

Policy Iteration

□ Initialize the Q functions with $Q_0(s, a) = R(s, a)$

□ Iterate as

$$Q^{new}(s, a) \leftarrow R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a' \in A} Q^{old}(s', a')$$

□ May be shown that for a large number of steps this iterative solution converges to $Q^\pi(s, a)$ for the optimal policy $\pi(a; s)$

Practical Limitations of MDP Policy Learning

$$Q^{new}(s, a) \leftarrow R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a' \in A} Q^{old}(s', a')$$

- ❑ Assumes we have access to a model of the environment, given by $P(s, a, s')$
- ❑ Typically we do not have these model parameters
- ❑ Possible strategy: Estimate $P(s, a, s')$ based on experience with environment
- ❑ Challenge: May require a lot of such experience to get good estimates

Appendix 2

Alternative Derivation of Q Learning

Learn the Policy as we Experience Environment

$$Q^{new}(s, a) \leftarrow R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a' \in A} Q^{old}(s', a')$$

Learn the Policy as we Experience Environment

$$Q^{new}(s, a) \leftarrow R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a' \in A} Q^{old}(s', a')$$

□ Based on the definition of $R(s, a)$, this is equivalent to

$$Q^{new}(s, a) \leftarrow \sum_{s' \in S} P(s, a, s') [r(s, a, s') + \gamma \max_{a' \in A} Q^{old}(s', a')]$$

Learn the Policy as we Experience Environment

$$Q^{new}(s, a) \leftarrow R(s, a) + \gamma \sum_{s' \in S} P(s, a, s') \max_{a' \in A} Q^{old}(s', a')$$

□ Based on the definition of $R(s, a)$ this is equivalent to

$$Q^{new}(s, a) \leftarrow \sum_{s' \in S} P(s, a, s') [r(s, a, s') + \gamma \max_{a' \in A} Q^{old}(s', a')]$$

□ By the definition of probability, as $K \rightarrow \infty$, this is equivalent to

$$Q^{new}(s, a) \leftarrow \frac{1}{T} \sum_{t=1}^T [r(s, a, s'_t) + \gamma \max_{a'} Q^{old}(s'_t, a')], \quad \text{with each } s'_t \sim P(s, a, s')$$

Learning from Experience

$$Q^{new}(s, a) \leftarrow \frac{1}{T} \sum_{t=1}^T [r(s, a, s'_t) + \gamma \max_{a'} Q^{old}(s'_t, a')], \quad \text{with each } s'_t \sim P(s, a, s')$$

- ❑ We don't know the underlying probability distribution $P(s, a, s')$
- ❑ But we can sample from it, from experience; don't need to know or estimate $P(s, a, s')$
- ❑ When in state s , try an action $a \in A$, and see what happens, *i.e.*, see what reward r is manifested
- ❑ Model learns to value, or *reinforce*, actions in a given state that are expected to be valuable

Naïve Q Learning

$$Q^{new}(s, a) \leftarrow \frac{1}{T} \sum_{t=1}^T [r(s, a, s'_t) + \gamma \max_{a'} Q^{old}(s'_t, a')] , \quad \text{with each } s'_t \sim P(s, a, s')$$

□ Suggests an update rule we implement “on the fly”, as we experience the environment

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + [r_t + \gamma \max_{a'} Q^{old}(s'_t, a')] - Q^{old}(s, a) , \quad \text{with } s'_t \sim P(s, a, s')$$

Naïve Q Learning

$$Q^{new}(s, a) \leftarrow \frac{1}{T} \sum_{t=1}^T [r(s, a, s'_t) + \gamma \max_{a'} Q^{old}(s'_t, a')] , \quad \text{with each } s'_t \sim P(s, a, s')$$

- ❑ Suggests an update rule we implement “on the fly”, as we experience the environment

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + [r_t + \gamma \max_{a'} Q^{old}(s'_t, a')] - Q^{old}(s, a) , \quad \text{with } s'_t \sim P(s, a, s')$$

- ❑ The $1/T$ factor doesn't affect policy (max operation), and therefore we ignore it for simplicity

Q-Learning

- ❑ The simple/naïve setup is

$$Q^{new}(s, a) \leftarrow Q^{old}(s, a) + [r_t + \gamma \max_{a'} Q^{old}(s'_t, a')]$$

- ❑ Introduce a “learning rate” $\alpha \in [0,1]$ and modify the update rule as

$$Q^{new}(s, a) \leftarrow (1 - \alpha) \cdot Q^{old}(s, a) + \alpha \cdot [r_t + \gamma \max_{a'} Q^{old}(s'_t, a')]$$

- ❑ This is called Q learning, and it allows one to learn a policy on the life, adaptively, as the environment is experienced

Appendix 3

Connecting Deep Q Learning to Conventional Q Learning

DQN Parameter Update Equation

- Recall that in deep Q learning the update equation of the model parameters is:

$$\theta^{new} \leftarrow \theta^{old} + \alpha \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta) |_{\theta=\theta^{old}}$$

- After implementing the parameter update, the Q function may be expressed as

$$Q(s_t, a_t; \theta^{new}) = Q(s_t, a_t; \theta^{old} + \alpha \cdot [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old})] \nabla_{\theta} Q(s_t, a_t; \theta) |_{\theta=\theta^{old}})$$

Taylor Series Expansion

$$Q(s_t, a_t; \theta^{new}) = Q(s_t, a_t; \theta^{old} + \alpha \cdot [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old})] \nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}})$$

□ Define

$$\Delta\theta = \alpha \cdot [r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old})] \nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}}$$

from which

$$Q(s, a; \theta^{new}) = Q(s, a; \theta^{old} + \Delta\theta)$$

□ First-order Taylor-series expansion, for small $\Delta\theta$:

$$Q(s, a; \theta^{old} + \Delta\theta) \approx Q(s, a; \theta^{old}) + \Delta\theta^T \nabla_{\theta} Q(s, a; \theta)|_{\theta=\theta^{old}}$$

Taylor Series Expansion

□ By the definition of $\Delta\theta$

$$Q(s_t, a_t; \theta^{old} + \Delta\theta) \approx Q(s_t, a_t; \theta^{old}) + \alpha \cdot \left[r_t + \gamma \cdot \max_a Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \|\nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}}\|_2^2$$

□ Define $\alpha' = \alpha \cdot \|\nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}}\|_2^2$, yielding approximately

$$Q(s_t, a_t; \theta^{new}) \leftarrow Q(s_t, a_t; \theta^{old}) \cdot (1 - \alpha') + \alpha' \cdot \left[r_t + \gamma \cdot \max_a Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right]$$

Connecting DQN to Q Learning

- The update rule for the model parameters in deep learning is

$$\theta^{new} \leftarrow \theta^{old} + \alpha \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}}$$

- When viewed from the perspective of the Q function this yields

$$Q(s_t, a_t; \theta^{new}) \leftarrow Q(s_t, a_t; \theta^{old}) \cdot (1 - \alpha') + \alpha' \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right]$$

$$\alpha' = \alpha \cdot \|\nabla_{\theta} Q(s_t, a_t; \theta)|_{\theta=\theta^{old}}\|_2^2$$

- Hence, the update rule in deep Q learning for the model parameters corresponds exactly to the update rule in conventional Q learning

Connecting DQN to Q Learning

- ❑ When doing updates in deep Q learning we do NOT actually implement

$$Q(s_t, a_t; \theta^{new}) \leftarrow Q(s_t, a_t; \theta^{old}) \cdot (1 - \alpha') + \alpha' \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right]$$

because now the Q function is not simply a table, as in original Q learning

- ❑ We implement

$$\theta^{new} \leftarrow \theta^{old} + \alpha \cdot \left[r_t + \gamma \cdot \max_{a'} Q(s'_{t+1}, a'; \theta^{old}) - Q(s_t, a_t; \theta^{old}) \right] \nabla_{\theta} Q(s_t, a_t; \theta) |_{\theta=\theta^{old}}$$

meaning we update the functional model $Q(s, a; \theta)$ for the Q function